

Gen. idea

- take two images, produce a new image that reflects the content of one but the artistic style of the other.
- loss function that matches the style & content of each resp. image in the feature space of the deep network, & then performing gradient descent on the pixels of the image itself.
- VGG19, AlexNet → feature extractor
 - ↳ trained on ImageNet

Loss

- weighted sum of 3 terms: content loss + style loss + total variation loss.
- penalize deviations from the content of the content image & deviations from the style of the style image.
- perform grad. descent not on the parameters of the model, but on the pixel values of our original image.

→ Content loss

- feature map of gene. image differs from feature map of source image.
- content repr / feature map of ^{one} layer: A ($1 \times C \times H \times W$). (reshape)
- F : feat. map of current image
- P : feat. map of content source image.
- $L_c = w_c \sum_{ij} (F_{ij} - P_{ij})^2$

→ Style loss

- Compute gram matrix G : correlations betⁿ values in each channel of the feature map.
- We want activation statistics of our generated image to match the activation stats of our style image.

$$G_{ij} = \sum_k F_{ik} F_{jk}$$

↳ Gram matrix from the feature map of the current image.

→ A : Gram matrix from the feature map of the ^{Source} style image.

$$L_s^l = w_s \sum_{ij} (G_{ij} - A_{ij})^2$$

$$L_s = \sum_{\text{set of layers}} L_s^l$$

→ Total variation regul^r (used in perceptual losses for real-time style transf.)

spatial smoothness

- that penalises the wiggles / total variation in pixel values.

- sum of squares of difference in pixel values for all pairs of pixels that are next to each other

$$L_{VR} = w_{VR} \left(\sum_{c=1}^3 \sum_{i=1}^{H-1} \sum_{j=1}^W (x_{i+1,c} - x_{ij,c})^2 + \sum_{c=1}^3 \sum_{i=1}^H \sum_{j=1}^{W-1} (x_{i,j+1,c} - x_{ij,c})^2 \right)$$

Models:

- VGG-16, 19 pretrained
- Our own architecture (^{RESNEXT} ResNet, GoogleNet, good stem & head)
- Model trained on Assignment 2

Dataset:

- ↳ ImageNet

Dash

- ↳ fast-dash by Kedar.

